


Local Search and Optimization

Outline

- Local search techniques and optimization
 - Hill-climbing

Local search and optimization


- Previously: systematic exploration of search space.
 - Backtrack search
 - Can solve n-queen problems for $n = 200$
- Different algorithms can be used
 - Local search
 - Can solve n-queen for $n = 1,000,000$



Local search and optimization

- Local search
 - Keep track of single current state
 - Move only to neighboring states
 - Ignore paths
- Advantages:
 - Use very little memory
 - Can often find reasonable solutions in large or infinite (continuous) state spaces.
- "Pure optimization" problems
 - All states have an objective function
 - Goal is to find state with max (or min) objective value
 - Does not quite fit into CSP formulation
 - Local search can do quite well on these problems.

"Landscape" of search for max value



Hill-climbing search

```

function HILL-CLIMBING( problem ) return a state that is a local maximum
input: problem, a problem
local variables: current, a node.
                  neighbor, a node.

current  $\leftarrow$  MAKE-NODE(INITIAL-STATE[problem])
loop do
  neighbor  $\leftarrow$  a highest valued successor of current
  if VALUE [neighbor]  $\leq$  VALUE [current] then return STATE [current]
  current  $\leftarrow$  neighbor
  
```


This version of HILL-CLIMBING found local maximum.

Hill-climbing search

- "a loop that continuously moves in the direction of increasing value"
 - terminates when a peak is reached
 - Aka greedy local search
- Value can be either
 - Objective function value
 - Heuristic function value (minimized)
- Hill climbing does not look ahead of the immediate neighbors of the current state.
- Can randomly choose among the set of best successors, if multiple have the best value
- Characterized as "trying to find the top of Mount Everest while in a thick fog"

Hill climbing and local maxima


- When local maxima exist, hill climbing is suboptimal
- Simple (often effective) solution
 - Multiple random restarts



Hill-climbing example

- 8-queens problem, complete-state formulation
 - All 8 queens on the board in some configuration
- Successor function:
 - move a single queen to another square in the same column.
- Example of a heuristic function $h(n)$:
 - the number of pairs of queens that are attacking each other (directly or indirectly)
 - (so we want to minimize this)

Hill-climbing example




$$(c_1 c_2 c_3 c_4 c_5 c_6 c_7 c_8) = (5 6 7 4 5 6 7 6)$$

Current state: $h=17$


Shown is the h-value for each possible successor in each column

A local minimum for 8-queens



A local minimum in the 8-queens state space ($h=1$)

Other drawbacks



- Ridge = sequence of local maxima difficult for greedy algorithms to navigate
- Plateau = an area of the state space where the evaluation function is flat.

Performance of hill-climbing on 8-queens

- Randomly generated 8-queens starting states...
- 14% the time it solves the problem
- 86% of the time it get stuck at a local minimum
- However...
 - Takes only 4 steps on average when it succeeds
 - And 3 on average when it gets stuck
 - (for a state space with ~17 million states)

Possible solution...sideways moves

- If no downhill (uphill) moves, allow sideways moves in hope that algorithm can escape
 - Need to place a limit on the possible number of sideways moves to avoid infinite loops
- For 8-queens
 - Now allow sideways moves with a limit of 100
 - Raises percentage of problem instances solved from 14 to 94%
 - However....
 - 21 steps for every successful solution
 - 64 for each failure

Hill-climbing variations

- Stochastic hill-climbing
 - Random selection among the uphill moves.
 - The selection probability can vary with the steepness of the uphill move.
- First-choice hill-climbing
 - stochastic hill climbing by generating successors randomly until a better one is found
 - Useful when there are a very large number of successors
- Random-restart hill-climbing
 - Tries to avoid getting stuck in local maxima.

Hill-climbing with random restarts

- Different variations
 - For each restart: run until termination v. run for a fixed time
 - Run a fixed number of restarts or run indefinitely
- Analysis
 - Say each search has probability p of success
 - E.g., for 8-queens, $p = 0.14$ with no sideways moves
 - Expected number of restarts?
 - Expected number of steps taken?

Search using Simulated Annealing

- Simulated Annealing = hill-climbing with non-deterministic search
- Basic ideas:
 - like hill-climbing identify the quality of the local improvements
 - instead of picking the best move, pick one randomly
 - say the change in objective function is δ
 - if δ is positive, then move to that state
 - otherwise:
 - move to this state with probability proportional to δ
 - thus: worse moves (very large negative δ) are executed less often
 - however, there is always a chance of escaping from local maxima
 - over time, make it less likely to accept locally bad moves
 - (Can also make the size of the move random as well, i.e., allow "large" steps in state space)

Physical Interpretation of Simulated Annealing

- Annealing = physical process of cooling a liquid or metal until particles achieve a certain frozen crystal state
 - simulated annealing:
 - free variables are like particles
 - seek "low energy" (high quality) configuration
 - get this by slowly reducing temperature T , which particles move around randomly

Simulated annealing

```

function SIMULATED-ANNEALING( problem, schedule) return a solution state
  input: problem, a problem
           schedule, a mapping from time to temperature
  local variables: current, a node.
                   next, a node.
                   T, a "temperature" controlling the probability of downward steps

  current  $\leftarrow$  MAKE-NODE(INITIAL-STATE[problem])
  for t  $\leftarrow$  1 to  $\infty$  do
    T  $\leftarrow$  schedule[t]
    if T = 0 then return current
    next  $\leftarrow$  a randomly selected successor of current
     $\Delta E \leftarrow$  VALUE[next] - VALUE[current]
    if  $\Delta E > 0$  then current  $\leftarrow$  next
    else current  $\leftarrow$  next only with probability  $e^{\Delta E / T}$ 

```

More Details on Simulated Annealing

- Lets say there are 3 moves available, with changes in the objective function of $d1 = -0.1$, $d2 = 0.5$, $d3 = -5$. (Let $T = 1$).
- pick a move randomly:
 - if $d2$ is picked, move there.
 - if $d1$ or $d3$ are picked, probability of move = $\exp(d/T)$
 - move 1: $\text{prob1} = \exp(-0.1) = 0.9$,
 - i.e., 90% of the time we will accept this move
 - move 3: $\text{prob3} = \exp(-5) = 0.05$
 - i.e., 5% of the time we will accept this move
- T = "temperature" parameter
 - high $T \Rightarrow$ probability of "locally bad" move is higher
 - low $T \Rightarrow$ probability of "locally bad" move is lower
 - typically, T is decreased as the algorithm runs longer
 - i.e., there is a "temperature schedule"

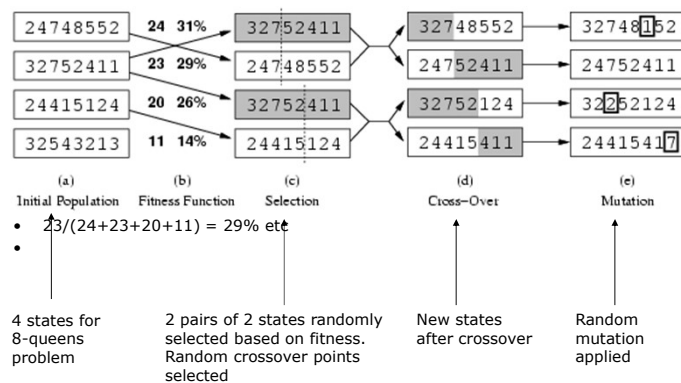
Simulated Annealing in Practice

- method proposed in 1983 by IBM researchers for solving VLSI layout problems (Kirkpatrick et al, *Science*, 220:671-680, 1983).
 - theoretically will always find the global optimum (the best solution)
- useful for some problems, but can be very slow
 - slowness comes about because T must be decreased very gradually to retain optimality
- In practice how do we decide the rate at which to decrease T ? (this is a practical problem with this method)


Genetic algorithms

- Different approach to other search algorithms
 - A successor state is generated by combining two parent states
- A state is represented as a string over a finite alphabet (e.g. binary)
 - 8-queens
 - State = position of 8 queens each in a column
 - $\Rightarrow 8 \times \log(8)$ bits = 24 bits (for binary representation)
- Start with k randomly generated states (population)
- Evaluation function (fitness function).
 - Higher values for better states.
 - Opposite to heuristic function, e.g., # non-attacking pairs in 8-queens
- Produce the next generation of states by "simulated evolution"
 - Random selection
 - Crossover
 - Random mutation

Genetic algorithms



Genetic algorithms



Has the effect of "jumping" to a completely different new part of the search space (quite non-local)

Genetic algorithm pseudocode

```

function GENETIC_ALGORITHM( population, FITNESS-FN) return an individual
  input: population, a set of individuals
         FITNESS-FN, a function which determines the quality of the individual
  repeat
    new_population  $\leftarrow$  empty set
    loop for i from 1 to SIZE(population) do
      x  $\leftarrow$  RANDOM_SELECTION(population, FITNESS_FN)
      y  $\leftarrow$  RANDOM_SELECTION(population, FITNESS_FN)
      child  $\leftarrow$  REPRODUCE(x,y)
      if (small random probability) then child  $\leftarrow$  MUTATE(child )
      add child to new_population
    population  $\leftarrow$  new_population
  until some individual is fit enough or enough time has elapsed
  return the best individual

```

Comments on genetic algorithms

- Positive points
 - Random exploration can find solutions that local search can't
 - (via crossover primarily)
 - Appealing connection to human evolution
 - E.g., see related area of genetic programming
- Negative points
 - Large number of "tunable" parameters
 - Difficult to replicate performance from one problem to another
 - Lack of good empirical studies comparing to simpler methods
 - Useful on some (small?) set of problems but no convincing evidence that GAs are better than hill-climbing w/random restarts in general