

Solving problems by searching

1

Outline

- Problem-solving agents
- Problem types
- Problem formulation
- Example problems
- Basic search algorithms

2

Problem-solving agents

```

function SIMPLE-PROBLEM-SOLVING-AGENT(percept) returns an action
  static: seq, an action sequence, initially empty
           state, some description of the current world state
           goal, a goal, initially null
           problem, a problem formulation

  state ← UPDATE-STATE(state, percept)
  if seq is empty then do
    goal ← FORMULATE-GOAL(state)
    problem ← FORMULATE-PROBLEM(state, goal)
    seq ← SEARCH(problem)
  action ← FIRST(seq)
  seq ← REST(seq)
  return action

```

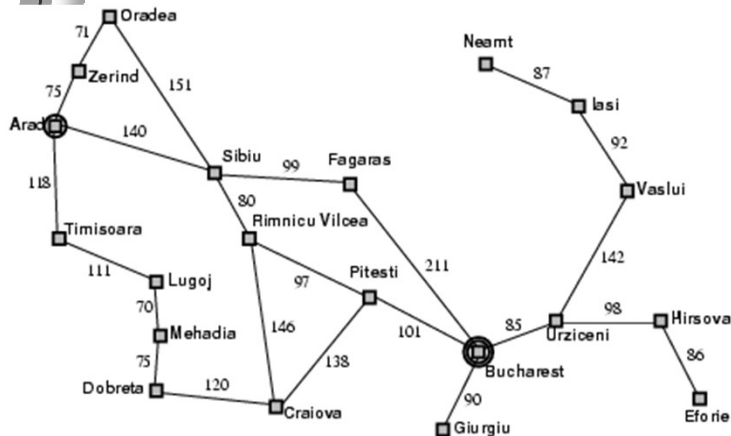
3

Example: Romania

- On holiday in Romania; currently in Arad.
- Flight leaves tomorrow from Bucharest
 - be in Bucharest
- Formulate problem:
 - states: various cities
 - actions: drive between cities
- Find solution:
 - sequence of cities, e.g., Arad, Sibiu, Fagaras, Bucharest
- Formulate goal

4

Example: Romania



5

Problem types

- Deterministic, fully observable → single-state problem
 - Agent knows exactly which state it will be in; solution is a sequence
- Non-observable → sensorless problem (conformant problem)
 - Agent may have no idea where it is; solution is a sequence
- Nondeterministic and/or partially observable → contingency problem
 - percepts provide new information about current state
 - often interleave} search, execution
- Unknown state space → exploration problem

6

Single-state problem formulation

A problem is defined by four items:

1. initial state e.g., "at Arad"
 - e.g., $S(\text{Arad}) = \{ \langle \text{Arad} \rightarrow \text{Zerind}, \text{Zerind} \rangle, \dots \}$
 2. goal test, can be
 - explicit, e.g., $x = \text{"at Bucharest"}$
 - implicit, e.g., $\text{Checkmate}(x)$
 3. path cost (additive)
 - e.g., sum of distances, number of actions executed, etc.
 - $c(x, a, y)$ is the step cost, assumed to be ≥ 0
- A solution is a sequence of actions leading from the initial state to a goal state

actions or successor function $S(x) = \text{set of action-state pairs}$


7

Selecting a state space

- Real world is absurdly complex
 - state space must be abstracted for problem solving
- (Abstract) state = set of real states
 - e.g., "Arad → Zerind" represents a complex set of possible routes, detours, rest stops, etc.
- For guaranteed realizability, any real state "in Arad" must get to some real state "in Zerind"
 - set of real paths that are solutions in the real world
- Each abstract action should be "easier" than the original problem
- (Abstract) action = complex combination of real actions

8


Vacuum world state space graph



- states? integer dirt and robot location
- actions? *Left, Right, Suck*
- goal test? no dirt at all locations
- path cost? 1 per action

9


Example: The 8-puzzle



- states?
- actions?
- goal test?
- path cost?

10


Example: The 8-puzzle



- states? locations of tiles
 - actions? move blank left, right, up, down
 - goal test? = goal state (given)
 - path cost? 1 per move
- [Note: optimal solution of n -Puzzle family is NP-hard]

11

Example: robotic assembly



- states?: real-valued coordinates of robot joint angles parts of the object to be assembled
- actions?: continuous motions of robot joints
- goal test?: complete assembly
- path cost?: time to execute

12

Tree search algorithms

- Basic idea:
 - offline, simulated exploration of state space by generating successors of already-explored states (a.k.a. \sim expanding states)


```

function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree

```


13

Tree search example




14

Tree search example



15

Tree search example



16

Implementation: general tree search

```

function TREE-SEARCH(problem, fringe) returns a solution, or failure
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)
    fringe ← INSERTALL(EXPAND(node, problem), fringe)


function EXPAND(node, problem) returns a set of nodes
  successors ← the empty set
  for each action, result in SUCCESSOR-FN[problem](STATE[node]) do
    s ← a new NODE
    PARENT-NODE[s] ← node; ACTION[s] ← action; STATE[s] ← result
    PATH-COST[s] ← PATH-COST[node] + STEP-COST(node, action, s)
    DEPTH[s] ← DEPTH[node] + 1
    add s to successors
  return successors

```

17

Implementation: states vs. nodes

- A state is a (representation of) a physical configuration
- A node is a data structure constituting part of a search tree includes state, parent node, action, path cost $g(x)$, depth



- The `Expand` function creates new nodes, filling in the various fields and using the `SuccessorFn` of the problem to create the corresponding states.

18

Search strategies

- A search strategy is defined by picking the order of node expansion
- Strategies are evaluated along the following dimensions:
 - completeness: does it always find a solution if one exists?
 - time complexity: number of nodes generated
 - space complexity: maximum number of nodes in memory
 - optimality: does it always find a least-cost solution?
- Time and space complexity are measured in terms of
 - b : maximum branching factor of the search tree
 - d : depth of the least-cost solution
 - m : maximum depth of the state space (may be ∞)

19


Uninformed search strategies

- Uninformed search strategies use only the information available in the problem definition
- Breadth-first search
- Uniform-cost search
- Depth-first search
- Depth-limited search
- Iterative deepening search

20

Breadth-first search


- Expand shallowest unexpanded node
 - fringe* is a FIFO queue, i.e., new successors go at end
- Implementation:



21

Breadth-first search


- Expand shallowest unexpanded node
 - fringe* is a FIFO queue, i.e., new successors go at end
- Implementation:



22

Breadth-first search


- Expand shallowest unexpanded node
 - fringe* is a FIFO queue, i.e., new successors go at end
- Implementation:



23

Breadth-first search

- Expand shallowest unexpanded node
 - fringe* is a FIFO queue, i.e., new successors go at end
- Implementation:



24

Properties of breadth-first search

- Complete? Yes (if b is finite)
- Time? $1+b+b^2+b^3+\dots +b^d + b(b^d-1) = O(b^{d+1})$
- Space? $O(b^{d+1})$ (keeps every node in memory)
- Optimal? Yes (if cost = 1 per step)
- Space is the bigger problem (more than time)

25

Uniform-cost search

- Expand least-cost unexpanded node
fringe = queue ordered by path cost
- Equivalent to breadth-first if step costs all equal
- Space? # of nodes with $g \leq \text{cost of optimal solution}$, $O(b^{\text{ceiling}(C^*/\epsilon)})$
- Optimal? Yes – nodes expanded in increasing order of $g(n)$
- Complete? Yes, if step cost $\geq \epsilon$
- Time? # of nodes with $g \leq \text{cost of optimal solution}$, $O(b^{\text{ceiling}(C^*/\epsilon)})$ where C^* is the cost of the optimal solution

26

Depth-first search

- Expand deepest unexpanded node
 - *fringe* = LIFO queue, i.e., put successors at front
- Implementation:



27

Depth-first search

- Expand deepest unexpanded node
 - *fringe* = LIFO queue, i.e., put successors at front
- Implementation:



28

Depth-first search

- Expand deepest unexpanded node
 - fringe* = LIFO queue, i.e., put successors at front
- Implementation:



29

Depth-first search

- Expand deepest unexpanded node
 - fringe* = LIFO queue, i.e., put successors at front
- Implementation:



30

Depth-first search

- Expand deepest unexpanded node
 - fringe* = LIFO queue, i.e., put successors at front
- Implementation:



31

Depth-first search

- Expand deepest unexpanded node
 - fringe* = LIFO queue, i.e., put successors at front
- Implementation:



32

Depth-first search

- Expand deepest unexpanded node
 - fringe* = LIFO queue, i.e., put successors at front
- Implementation:



33

Depth-first search

- Expand deepest unexpanded node
 - fringe* = LIFO queue, i.e., put successors at front
- Implementation:



34

Depth-first search


- Expand deepest unexpanded node
 - fringe* = LIFO queue, i.e., put successors at front
- Implementation:



35

Depth-first search


- Expand deepest unexpanded node
 - fringe* = LIFO queue, i.e., put successors at front
- Implementation:



36

Depth-first search


- Expand deepest unexpanded node
 - fringe* = LIFO queue, i.e., put successors at front
- Implementation:



37

Depth-first search

- Expand deepest unexpanded node
 - fringe* = LIFO queue, i.e., put successors at front
- Implementation:



38

Properties of depth-first search

- Complete? No: fails in infinite-depth spaces, spaces with loops
 - Modify to avoid repeated states along path
- Time? $O(b^m)$: terrible if m is much larger than d
 - but if solutions are dense, may be much faster than breadth-first
- Space? $O(bm)$, i.e., linear space!
- Optimal? No
 - complete in finite spaces

39

Depth-limited search

= depth-first search with depth limit l ,
i.e., nodes at depth l have no successors

- Recursive implementation:

```
function DEPTH-LIMITED-SEARCH(problem, limit) returns soln/fail/cutoff
  RECURSIVE-DLS(MAKE-NODE(INITIAL-STATE[problem]), problem, limit)
function RECURSIVE-DLS(node, problem, limit) returns soln/fail/cutoff
  cutoff-occurred? ← false
  if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)
  else if DEPTH[node] = limit then return cutoff
  else for each successor in EXPAND(node, problem) do
    result ← RECURSIVE-DLS(successor, problem, limit)
    if result = cutoff then cutoff-occurred? ← true
    else if result ≠ failure then return result
  if cutoff-occurred? then return cutoff else return failure
```

40

Iterative deepening search

function ITERATIVE-DEEPENING-SEARCH(*problem*) returns a solution, or failure

inputs: *problem*, a problem

for *depth* $\leftarrow 0$ to ∞ do

result \leftarrow DEPTH-LIMITED-SEARCH(*problem*, *depth*)

 if *result* \neq cutoff then return *result*

41

Iterative deepening search /=0

Limit = 0



42

Iterative deepening search /=1


Limit = 1



43


Iterative deepening search /=2

Limit = 2



44

Iterative deepening search $d=3$



45

Iterative deepening search

- Number of nodes generated in a depth-limited search to depth d with branching factor b :

$$N_{DLS} = b^0 + b^1 + b^2 + \dots + b^{d-2} + b^{d-1} + b^d$$

- Number of nodes generated in an iterative deepening search to depth d with branching factor b :

$$N_{IDS} = (d+1)b^0 + d b^1 + (d-1)b^2 + \dots + 3b^{d-2} + 2b^{d-1} + 1b^d$$

For $b = 10, d = 5$,

$$N_{DLS} = 1 + 10 + 100 + 1,000 + 10,000 + 100,000 = 111,111$$

- Overhead = $(123,456 - 111,111)/111,111 = 11\%$
 - $NIDS = 6 + 50 + 400 + 3,000 + 20,000 + 100,000 = 123,456$

46

Properties of iterative deepening search

- Complete? Yes
- Time? $(d+1)b^0 + d b^1 + (d-1)b^2 + \dots + b^d = O(bd)$
- Space? $O(bd)$
- Optimal? Yes, if step cost = 1

47

Summary of algorithms

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening
Complete?	Yes	Yes	No	No	Yes
Time	$O(b^{d+1})$	$O(b^{lC^*/\epsilon})$	$O(b^m)$	$O(b^l)$	$O(b^d)$
Space	$O(b^{d+1})$	$O(b^{lC^*/\epsilon})$	$O(bm)$	$O(bl)$	$O(bd)$
Optimal?	Yes	Yes	No	No	Yes

48

Graph search

```

function GRAPH-SEARCH(problem, fringe) returns a solution, or failure
  closed ← an empty set
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)
    if STATE[node] is not in closed then
      add STATE[node] to closed
      fringe ← INSERTALL(EXPAND(node, problem), fringe)

```

49

Summary

- Problem formulation usually requires abstracting away real-world details to define a state space that can feasibly be explored
- Variety of uninformed search strategies
- Iterative deepening search uses only linear space and not much more time than other uninformed algorithms

50